

Combining Gradient-Based With Evolutionary Online Learning: An Introduction to Learning Classifier Systems

Martin V. Butz

Department of Psychology, Cognitive Psychology III
University of Würzburg, Röntgenring 11, 97070, Würzburg, Germany
butz@psychologie.uni-wuerzburg.de

Abstract

Learning Classifier Systems (LCSs), introduced by John H. Holland in the 1970s, are rule-based evolutionary online learning systems that combine gradient-based rule evaluation with evolutionary-based rule structuring techniques. Since the introduction of the accuracy-based XCS classifier system by Stewart W. Wilson in 1995, LCSs showed to be flexible, online learning methods that are applicable to datamining, reinforcement learning, and function approximation problems. Comparisons showed that performance is competitive with state-of-the art machine learning algorithms, but the learning algorithms applied are usually more flexible and highly adaptive. Moreover, problem knowledge can be extracted easily. This tutorial provides a gentle introduction to LCSs and their general functioning. It then gives further details on the XCS classifier system and highlights various successful applications. In conclusion, promising future directions of LCS research and applications are discussed.

1. Introduction

This tutorial introduces Learning Classifier Systems (LCSs) [6, 19, 21] from a hybrid intelligent systems perspective. LCSs combine machine learning techniques with evolutionary learning techniques to evolve distributed problem solutions. These solutions are represented by a set of rules, the so-called *population of classifiers*. Each classifier is locally applicable in a problem subspace, specified in a rule condition, and provides a problem solution classification or prediction. Classifiers compete within their specified problem space with other, overlapping classifiers for reproductions. The classifiers that win the competition represent more *suitable* problem solutions, as specified in the chosen fitness approach.

Different classifier implementations modify the notion of

fitness, the machine learning techniques used for the generation of prediction and fitness estimates, the representation of conditions and predictions, as well as the evolutionary learning component. Two general types of LCSs can be distinguished: Michigan-style and Pittsburgh-style LCSs. Michigan-style LCSs were originally introduced as *Cognitive Systems* [19, 21], simulating some animal-like behavior. In general, Michigan-style LCSs evolve iteratively online one global problem solution in their population of classifiers. Each classifier represents a solution in a problem subspace and competes with other classifiers for reproductions. Pittsburgh-style approaches evolve multiple problem solutions so that complete problem solutions, represented by a set of classifiers, compete for reproductions [1, 15, 26]. Since Pittsburgh-style LCSs are very similar to genetic algorithms [17] and usually do not integrate gradient-based update mechanisms, we focus on Michigan-style LCSs in the remainder of this tutorial.

We further focus on the currently most applied Michigan-style LCS: the XCS Classifier System, which was introduced by Stewart W. Wilson in 1995 [35]. XCS has been shown both to be flexibly applicable in various problem domains, including classification problems (datamining), reinforcement learning problems, and function approximation problems. and may be the most applied LCS today.

First, we now provide some historical background on LCS research. Next, we introduce LCSs from a very general perspective specifying the general framework and basic components. We then focus on the XCS classifier system providing further system details as well as exemplary applications. A short outlook concludes the paper.

2. Historical Remarks

Inspired by stimulus-response theories, behaviorism, and other psychological principles of learning and behavior at the time, John H. Holland proposed the basic LCS frame-

work as a *cognitive system* model [19, 21]. Holland designed his cognitive system to evolve a set of production rules. These rules converted (sensory or perceptual) input into useful (motor or internal state) output. Internal states were represented in a *message list*, which could situate the system in its current environmental context. Later, Holland proposed the infamous *bucket-brigade* algorithm [20], which predated later temporal difference learning approaches such as TD(λ) or SARSA [30].

Consequently, most of the early LCS implementations focused on the cognitive system aspects. In the first classifier system implementation, Holland and Reitman [21] designed an LCS that evolved goal-directed, stimulus-response-based behavior, which satisfied multiple needs, or motivations, represented in internal resource reservoirs. Booker [5] extended Holland's approach by experimenting with an agent that needs to avoid aversive stimuli and reach attractive stimuli. Wilson confirmed the potential of LCSs to simulate artificial animals, which he termed *animats* [32, 33].

However, LCSs were not only applied as cognitive systems. Goldberg [16] successfully applied an LCS to the control of a simulated pipeline system, confirming that LCSs are valuable learning systems for real-world applications as well. Despite the variety and promising initial results, the hybrid nature of the learning system in LCSs made it hard to analyze and comprehend the interaction of the mechanisms at work in the LCS implementations at the time. In their "*Critical review of classifier systems*", Wilson and Goldberg [38] pointed out several unresolved questions that made it hard to design working LCS systems without the need of many trials and errors. The most important challenges raised were: (1) the learning and maintenance of reward chains by the bucket-brigade algorithm; (2) the rule fitness approaches, which often obstructed generalization, enabled overgeneralization, or prevented the formation of useful hierarchical structures; (3) rule syntax, which prohibited effective feature processing as well as the effective coverage of continuous problem spaces or large binary spaces. Furthermore, it was suggested that planning and lookahead mechanisms, representations of expectations, implementations of short-term memory, and population sizing equations need to be developed further and understood in more detail.

Unfortunately—and probably due to the lack of solid theoretical knowledge on any of the aspects—LCSs did not receive much research attention for several years. One of the few researchers that continued to work with LCSs was Stewart Wilson who then heralded an LCS renaissance by designing two of the most influential LCS systems to date: (1) the zeroth level classifier system ZCS [34] and (2) the accuracy-based classifier system XCS [35]. Both classifier systems overcome many of the previously en-

countered challenges. The credit assignment mechanism in ZCS and XCS is directly related to the well-understood Q-learning algorithm [31]. This enabled effective rule reward estimation and propagation. Overgeneralization problems were overcome by proper fitness sharing techniques or an accuracy-based fitness approach. Meanwhile, effective generalization was achieved by a niche reproduction combined with population-wide deletion, as stated in Wilson's generalization hypothesis [35].

Theoretical analyses and applications later confirmed the great potential of XCS and its offspring systems. Investigations showed machine learning competitiveness in datamining applications, in the reinforcement learning domain, as well as in function approximation problems [2, 3, 7, 8, 14, 23, 36, 37]. However, before we delve into the XCS system, we first give a general overview of LCSs and their interacting representations and learning algorithms.

3. Learning Classifier Systems

Michigan-style LCSs [19, 21, 6] are rule-based evolutionary online learning systems [8]. Generally, an LCS consists of (1) a set of rules, that is, a *population of classifiers*, (2) a rule evaluation mechanism, and (3) a rule evolution mechanism.

System knowledge is represented in the population, which encodes the (evolving) problem solution. Each classifier consists usually out of (i) a condition part, (ii) an action part, and (iii) a prediction part. The condition part specifies in which problem subspace the classifier is applicable. The action part specifies the proposed classification or action. The prediction part specifies a function value or reward prediction, dependent on the problem type the LCS is applied to.

The prediction part and the correlated rule evaluation mechanisms are usually realized by gradient-based (often linear) approximation approaches. The original prediction updates in ZCS and XCS [34, 35] were derived from reinforcement learning (RL) [30] techniques and effectively replaced the bucket-brigade algorithm. Rule evaluation, that is, the determination of rule fitness, is then usually derived from the prediction estimate. Two approaches have shown most useful: A fitness approach based on reward sharing, as done in ZCS [34, 22] or a fitness approach based on the *accuracy* of reward predictions [35, 22].

Based on the estimated utilities, an evolutionary computation approach [18, 17, 28] generates offspring classifiers and deletes less useful classifiers. Thus, based on the fitness estimates, the evolutionary approach searches for rule structures that can improve prediction and fitness even further.

Learning takes place iteratively online. Each iteration a problem instance is received and processed by the LCS. First, those classifiers in the population are identified that

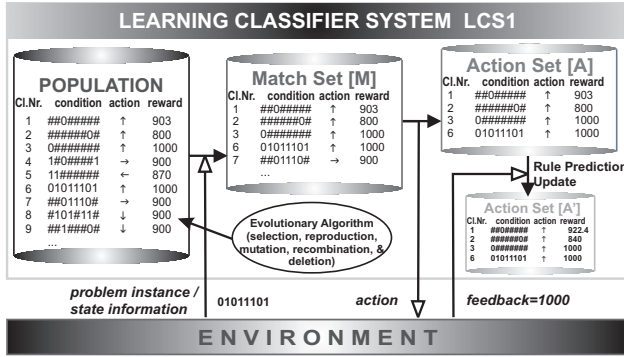


Figure 1. LCS iteratively interacts and learns from a problem / environment.

match the current given problem instance forming a *match set*. The classifiers in the match set essentially specify the current knowledge of the LCS about the problem instance. Consequently, the match set is used to decide on the classification, action, or prediction, dependent on the problem type and exploration strategy applied. If an action is chosen, an *action set* is formed that contains all classifiers in the match set that specify the chosen action. Once feedback is available about the decision made, predictions and fitness values of the classifiers in the action set are updated. Finally, the evolutionary component is applied. Usually, a steady-state genetic algorithm reproduces some classifiers and deletes others based on the fitness estimates. Figure 1 illustrates the learning process.

Thus, two learning mechanisms are at work. Gradient-based methods generate maximally suitable predictions and derive the fitness measure out of it. Dependent on the fitness measure used, classifiers are reproduced that yield high payoff, as done in the ZCS system [34], or also that generate accurate predictions, as done in the XCS system [35]. Thus, since the evolutionary algorithm relies on a fitness estimate generated by the chosen gradient-based method applied, it is essential that the LCS is designed in such a way that the gradient method identifies better classifiers as fast as possible so that the evolutionary component can lead the evolutionary process towards a suitable solution representation.

4. Problem Types

Generally, an LCS can be applied to any type of prediction problem imaginable. LCSs are designed to evolve solutions by partitioning the problem space into subspaces in which suitable predictions can be formed. Thus, LCSs are designed to detect those problem input features, or input subspaces, that are relevant to form effective or accurate predictions. We now discuss three distinct, albeit related,

problem types and the problem solution representation that an LCS would typically evolve in such a problem.

Classification Problems A classification problem is typically defined over a set of problem instances that consists of a set of features and a corresponding problem class. For example, in datamining, we intend to analyze datasets that contain various problem instances.

For example, the UCI repository [4] contains a mushroom dataset. Each instance in the dataset describes properties of a particular mushroom and partitions the mushrooms into edible and poisonous ones. The task for an LCS would then be to learn to be able to decide if a mushroom, given its properties, is edible or not. Classification problems can have various input features, such as nominal ones, real-valued ones, and binary ones and they may also contain more than two problem classes.

Comparisons of various LCSs with standard machine learning techniques have shown that the classification quality is comparable to those of state-of-the-art machine learning techniques [2, 3, 8]. Also binary classification problems, such as the well-studied multiplexer problem [8, 35, 36], are part of this problem class.

Reinforcement Learning Problems Besides classification problems, LCSs have been successfully applied to behavioral policy learning problems, typically describable by a Markov decision process [30]. In RL problems, feedback is provided only in the form of reward values, which needs to be propagated over the problem space to be able to decide on a proper action to execute. Thus, LCSs learn to propagate reward and evolve those classifiers that either predict the reward accurately or that predict the correct action accurately. While classification problems are usually single-step problems in which the successive problem inputs are independent of each other, in RL problems successive instances typically depend on each other and a solution can only be found if reward is *propagated* throughout the problem space.

A typical example of an RL problem is a path finding problem in a maze, in which reward is only provided at certain goal locations and the learning algorithm has to find the shortest path to the goal (or the closest goal) from all positions in the maze. LCS systems, and in particular the XCS classifier system, have shown robust performance in various maze tasks [11] but also in the mountain car problem [25] and blocks world problems [8, 10].

Function Approximation Problems Besides classification and RL problems, another class of problems can be described as function approximation problems. In these problems, the LCS learns the underlying function of a particular problem. That is, given certain input values, the LCS

predicts the output value(s). Often, LCSs evolve overlapping, piecewise linear function value approximations, as is the case for the XCS system [37, 7].

The wide variety of problems above suggests that there are multiple domains of potential applications of LCSs. We now proceed to specify the general structure of the XCS classifier system, its learning approach, and its solution representation.

5. The XCS Classifier System

The XCS classifier system [35, 36] is an LCS that evolves its classifiers by an accuracy-based fitness approach. XCS has been successfully applied to various problem types and may be the most widely used classifier system today. We now first give an overview over the system and next highlight some of the major successful applications.

5.1. XCS in a Nutshell

XCS's main differences to earlier LCS systems are (1) its accuracy-based fitness approach and (2) its niche based reproduction process. To accomplish this, several modifications in classifier representation and GA application were necessary.

Knowledge Representation As other LCSs, XCS represents its knowledge by a population of classifiers. The population size in XCS is usually set to a fixed upper bound. Each classifier contains the usual condition, action, and reward prediction parts. Additionally, though, each classifier contains a prediction error estimate and a fitness estimate, which reflects the accuracy (inverse error) of a classifier relative to overlapping, competing classifiers. Thus, fitness represents the relative accuracy of a classifier with respect to current competing classifiers (those in common match or action sets).

Evolutionary Component Consequently, the evolutionary component propagates classifiers with more accurate prediction values. Since the evolutionary component evolves classifier structure (condition and action parts), structures evolve that allow the generation of maximally accurate predictions. Dependent on the problem and the condition representation used, the problem space is consequently clustered for the generation of accurate predictions.

Unlike population-wide reproduction *and* deletion in previous LCSs, XCS reproduces in the current match or action set while it continuously deletes classifiers in the whole population. Thus, the evolutionary component searches in the current problem niche, defined by the given problem instance, for better problem sub-solutions, based on the

matching classifiers. In this way, evolution searches more locally in problem subspaces. Moreover, since deletion still takes place in the whole population, an intrinsic generalization pressure applies because the average classifier condition in the match set covers on average a larger problem subspace than the average classifier condition in the population, so that (without additional fitness influence) more general classifiers will be reproduced and less general classifiers will be deleted.

This combination leads to the desired learning process that generates maximally accurate and maximally general classifiers. Consequently, the solution representation as a whole becomes maximally accurate and maximally general and is represented by partially overlapping classifiers [12, 13, 8].

5.2. Capabilities

After highlighting the novel features of the XCS classifier system, we now proceed to give an overview over the major system capabilities. We essentially highlight that XCS, in its various forms, can accurately solve large, challenging classification problems, RL problems, as well as function approximation problems.

Classification Various studies have shown that XCS can solve very large and noisy classification problems. Maybe the most common test function in the LCS literature is the multiplexer problem [34, 35]. It has been shown that XCS is able to even solve the highly challenging 70-bit multiplexer [13]. To solve the problem successfully, it requires a population size of at least $20k$ classifiers and about 3 million learning iterations [8]. With population sizes less than $N = 20k$, performance sometimes does not reach complete accuracy but gets stuck in a sub-optimal, overgeneral solution representation. Bounding models that determine a minimum population size to ensure successful learning with high probability can be found elsewhere (cf. [8] and references therein). Besides the successful solution of the multiplexer problem and other binary classification problems, XCS has also been applied to datamining problems. In these cases, several studies have shown that XCS performs machine-learning competitively, that is, it classifies data similarly well when compared to other machine learning techniques [2, 3, 8].

Reinforcement Learning Besides classification problems, XCS has been applied successfully to various reinforcement learning problems. Particularly, various maze tasks have been studied including mazes with discrete and continuous representations as well as mazes that require the formation of long reward chains. Hereby, it is important to use a stabilization component in the gradient-based update

to ensure learning success [11]. In discrete domains, it has been shown that XCS is particularly noise robust as well as very efficient to identify prediction-relevant parameters easily handling (and efficiently identifying and consequently ignoring) even 90 irrelevant additional, randomly fluctuating binary input features [8]. Most recently, XCS for continuous inputs has been shown to be able to also solve the Mountain car problem [25] effectively.

Function Approximation In the above applications, XCS mostly learned constant prediction values. Recently, though, this representation has been expanded to linear representations and others [24, 25, 37]. In function approximation problems, XCSF (the 'F' stands for function approximation) has been shown to solve various challenging functions including superimposed sine and ridge functions with partially overlapping linear approximators. Recent investigations [14] have shown that XCSF outperforms self-organizing map approaches, such as the NeuralGAS algorithm [27], and performs comparably to statistical local learning approaches, such as constructive incremental learning by Schaal and Atkeson [29]. Meanwhile, however, XCS is a much more general learning approach that can be applied to more problems than those suitable for NeuralGAS nets or the constructive incremental learning approach. With respect to scalability, it has been shown that XCS can handle even seven dimensional real-valued input spaces, developing, for example, an accurate approximation surface for the oblique sinusoidal function $f(x_1, \dots, x_7) = \sin(4\pi(x_1 + \dots + x_7))$.

XCSF is now also available online as a Java implementation [9]. Figure 2 shows a screenshot of a typical final problem solution representation of XCSF in the specified oblique sinusoidal sine function in a three dimensional input space. It can be seen that XCSF detects the oblique orientation of the sine function and consequently partitions the problem space along the axis (0,0,0) to (1,1,1) to be able to form maximally accurate linear predictions with maximally general classifier conditions. Note that the representation scales the condition sizes down in order to make the space partitioning observable.

6. Conclusions

This LCS introduction has provided a glimpse at the functioning and current capabilities of LCSs and the XCS classifier system in particular. LCS form problem solutions represented by partially overlapping subsolutions. It learns such solutions by a combination of gradient-based subsolution approximation and evolutionary-based subsolution structuring. More recent LCS research has begun to understand the dynamics of the systems sufficiently well, generating scalable and robust LCSs.




Figure 2. Representation of the conditions of a typical, final, accurate classifier population in a three dimensional oblique sine function.

In comparison to other machine learning techniques, performance has been shown to be competitive and sometimes even superior. While being competitive, the LCS technique is much more broadly applicable than the systems compared with. Current research focuses on the optimization of classifier operators, further application studies, and more detailed rigorous system analysis. What remains to be shown is in which problem types the flexibility of the evolutionary algorithm in combination with efficient gradient-based approximation techniques is superior to other machine learning techniques. It seems that this will be particularly the case in problem domains in which either many local minima prevent other local, distributed gradient-based techniques to fail or in which gradient-based techniques cannot be applied at all since the gradient information can not be used directly to structure the problem space.

References

- [1] J. Bacardit. *Pittsburgh Genetic-Based Machine Learning in the Dame Mining Era: Representations, Generalization, and Run Time*. PhD thesis, Computer Engineering Department, University of Ramon Llull, Barcelona, Spain, 2004.
- [2] E. Bernadó, X. Llorà, and J. M. Garrell. XCS and GALE: A comparative study of two learning classifier systems and six other learning algorithms on classification tasks. In P. L. Lanzi, W. Stolzmann, and S. W. Wilson, editors, *Advances in Learning Classifier Systems (LNAI 2321)*, pages 115–132. Springer-Verlag, Berlin Heidelberg, 2002.
- [3] E. Bernadó-Mansilla and J. M. Garrell-Guiu. Accuracy-based learning classifier systems: Models, analysis, and ap-

- plications to classification tasks. *Evolutionary Computation*, 11:209–238, 2003.
- [4] C. Blake, E. Keogh, and C. Merz. UCI repository of machine learning databases, 1998. (<http://www.ics.uci.edu/mllearn/MLRepository.html>).
 - [5] L. B. Booker. *Intelligent Behavior as an Adaptation to the Task Environment*. PhD thesis, The University of Michigan, 1982.
 - [6] L. B. Booker. Classifier systems that learn internal world models. *Machine Learning*, 3:161–192, 1988.
 - [7] M. V. Butz. Kernel-based, ellipsoidal conditions in the real-valued XCS classifier system. *GECCO 2005: Genetic and Evolutionary Computation Conference*, pages 1835–1842, 2005.
 - [8] M. V. Butz. *Rule-Based Evolutionary Online Learning Systems: A Principled Approach to LCS Analysis and Design*. Springer-Verlag, Berlin Heidelberg, 2006.
 - [9] M. V. Butz. Documentation of XCSFJava 1.1 plus visualization. MEDAL Report 2007008, Missouri Estimation of Distribution Algorithms Laboratory, University of Missouri in St. Louis, MO, 2007.
 - [10] M. V. Butz and D. E. Goldberg. Generalized state values in an anticipatory learning classifier system. In M. V. Butz, O. Sigaud, and P. Gérard, editors, *Anticipatory Behavior in Adaptive Learning Systems: Foundations, Theories, and Systems*, pages 282–301. Springer-Verlag, Berlin Heidelberg, 2003.
 - [11] M. V. Butz, D. E. Goldberg, and P. L. Lanzi. Gradient descent methods in learning classifier systems: Improving XCS performance in multistep problems. *IEEE Transactions on Evolutionary Computation*, 9:452–473, 2005.
 - [12] M. V. Butz, D. E. Goldberg, and K. Tharakunnel. Analysis and improvement of fitness exploitation in XCS: Bounding models, tournament selection, and bilateral accuracy. *Evolutionary Computation*, 11:239–277, 2003.
 - [13] M. V. Butz, T. Kovacs, P. L. Lanzi, and S. W. Wilson. Toward a theory of generalization and learning in XCS. *IEEE Transactions on Evolutionary Computation*, 8:28–46, 2004.
 - [14] M. V. Butz, P. L. Lanzi, and S. W. Wilson. Function approximation with XCS: Hyperellipsoidal conditions, recursive least squares, and compaction. *IEEE Transactions on Evolutionary Computation*, in press.
 - [15] K. A. DeJong, W. M. Spears, and D. F. Gordon. Using genetic algorithms for concept learning. *Machine Learning*, 13(2/3):161–188, 1993.
 - [16] D. E. Goldberg. Computer-aided gas pipeline operation using genetic algorithms and rule learning. *Dissertation Abstracts International*, 44(10):3174B, 1983. Doctoral dissertation, University of Michigan.
 - [17] D. E. Goldberg. *Genetic Algorithms in Search, Optimization and Machine Learning*. Addison-Wesley, Reading, MA, 1989.
 - [18] J. H. Holland. *Adaptation in Natural and Artificial Systems*. University of Michigan Press, Ann Arbor, MI, 1975. second edition, 1992.
 - [19] J. H. Holland. Adaptation. In R. Rosen and F. Snell, editors, *Progress in theoretical biology*, volume 4, pages 263–293. Academic Press, New York, 1976.
 - [20] J. H. Holland. Properties of the bucket brigade algorithm. *Proceedings of an International Conference on Genetic Algorithms and their Applications*, pages 1–7, 1985.
 - [21] J. H. Holland and J. S. Reitman. Cognitive systems based on adaptive algorithms. In D. A. Waterman and F. Hayes-Roth, editors, *Pattern directed inference systems*, pages 313–329. Academic Press, New York, 1978.
 - [22] T. Kovacs. *Strength or Accuracy: Credit Assignment in Learning Classifier Systems*. Springer-Verlag, Berlin Heidelberg, 2003.
 - [23] P. L. Lanzi. An analysis of generalization in the XCS classifier system. *Evolutionary Computation*, 7(2):125–149, 1999.
 - [24] P. L. Lanzi, D. Loiacono, S. W. Wilson, and D. E. Goldberg. Extending XCSF beyond linear approximation. *GECCO 2005: Genetic and Evolutionary Computation Conference*, pages 1827–1834, 2005.
 - [25] P. L. Lanzi, D. Loiacono, S. W. Wilson, and D. E. Goldberg. Classifier prediction based on tile coding. *GECCO 2006: Genetic and Evolutionary Computation Conference*, pages 1497–1504, 2006.
 - [26] X. Llorà and J. M. Garrell. Knowledge independent data mining with fine-grained parallel evolutionary algorithms. *Proceedings of the Third Genetic and Evolutionary Computation Conference (GECCO-2001)*, pages 461–468, 2001.
 - [27] T. M. Martinetz, S. G. Berkovitsch, and K. J. Schulten. “Neural-gas” network for vector quantization and its application to time-series prediction. *IEEE Transactions on Neural Networks*, 4:558–569, 1993.
 - [28] I. Rechenberg. *Evolutionsstrategie Optimierung technischer Systeme nach Prinzipien der biologischen Evolution*. Friedrich Frommann Verlag, Stuttgart-Bad Cannstatt, 1973.
 - [29] S. Schaal and C. G. Atkeson. Constructive incremental learning from only local information. *Neural Computation*, 10:2047–2084, 1998.
 - [30] R. S. Sutton and A. G. Barto. *Reinforcement learning: An introduction*. MIT Press, Cambridge, MA, 1998.
 - [31] C. J. C. H. Watkins. *Learning from Delayed Rewards*. PhD thesis, King’s College, Cambridge, UK, 1989.
 - [32] S. W. Wilson. Knowledge growth in an artificial animal. *Proceedings of an International Conference on Genetic Algorithms and Their Applications*, pages 16–23, 1985.
 - [33] S. W. Wilson. Classifier systems and the animat problem. *Machine Learning*, 2:199–228, 1987.
 - [34] S. W. Wilson. ZCS: A zeroth level classifier system. *Evolutionary Computation*, 2:1–18, 1994.
 - [35] S. W. Wilson. Classifier fitness based on accuracy. *Evolutionary Computation*, 3(2):149–175, 1995.
 - [36] S. W. Wilson. Generalization in the XCS classifier system. *Genetic Programming 1998: Proceedings of the Third Annual Conference*, pages 665–674, 1998.
 - [37] S. W. Wilson. Classifiers that approximate functions. *Natural Computing*, 1:211–234, 2002.
 - [38] S. W. Wilson and D. E. Goldberg. A critical review of classifier systems. *Proceedings of the Third International Conference on Genetic Algorithms*, pages 244–255, 1989.